

Music Information Retrieval A3

James Davidson - V00812527

March 17th, 2017

All up-to-date code for this can be found at <https://github.com/jamesthomasdavidson/Music-Classifier>

Question 1

.

```

1  import numpy as np
2  from sklearn import svm, datasets, linear_model
3  from sklearn.externals.joblib import Memory
4  from sklearn.model_selection import train_test_split
5  from sklearn.metrics import confusion_matrix
6  from sklearn.neighbors import KNeighborsClassifier
7
8  def get_data():
9      return datasets.load_svmlight_file('genres3.libsvm')
10
11 def SVC():
12     X, y = get_data()
13     print('Support Vector Machine')
14     X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
15     classifier = svm.SVC(kernel='linear', C=.8)
16     y_pred = classifier.fit(X_train, y_train).predict(X_test)
17     print("Confusion matrix: \n" + str(confusion_matrix(y_test, y_pred)))
18     print("Accuracy: " + str(classifier.score(X,y)) + '\n\n')
19
20 def SGD():
21     X, y = get_data()
22     print('Stochastic Gradient Descent')
23     X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
24     classifier = linear_model.SGDClassifier()
25     y_pred = classifier.fit(X_train, y_train).predict(X_test)
26     print("Confusion matrix: \n" + str(confusion_matrix(y_test, y_pred)))
27     print("Accuracy: " + str(classifier.score(X,y)) + '\n\n')
28
29 def NN():
30     X, y = get_data()
31     print('Nearest Neighbours')
32     X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
33     classifier = KNeighborsClassifier(n_neighbors=2)
34     y_pred = classifier.fit(X_train, y_train).predict(X_test)
35     print("Confusion matrix: \n" + str(confusion_matrix(y_test, y_pred)))
36     print("Accuracy: " + str(classifier.score(X,y)) + '\n\n')
37
38 SVC()
39 SGD()
40 NN()

```

```

./mkcollection -c classical.mf -l classical ../genres/classical
./mkcollection -c rock.mf -l rock ../genres/rock
./mkcollection -c hiphop.mf -l hiphop ../genres/hiphop
cat cl.mf hi.mf ro.mf > genres3.mf
bextract -sv genres3.mf -w genres3.arff

```

Using Weka:

=====ZeroR=====

Correctly Classified Instances	100	33.3333 %
Incorrectly Classified Instances	200	66.6667 %
Kappa statistic	0	
Mean absolute error	0.4444	
Root mean squared error	0.4714	
Relative absolute error	100 %	
Root relative squared error	100 %	
Total Number of Instances	300	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	1.000	1.000	0.333	1.000	0.500	0.000	0.500	0.333	classical
	0.000	0.000	0.000	0.000	0.000	0.000	0.500	0.333	hiphop
	0.000	0.000	0.000	0.000	0.000	0.000	0.500	0.333	rock
Weighted Avg.	0.333	0.333	0.111	0.333	0.167	0.000	0.500	0.333	

=== Confusion Matrix ===

a b c <-- classified as
100 0 0 | a = classical
100 0 0 | b = hiphop
100 0 0 | c = rock

=====NaiveBayesSimple=====

Correctly Classified Instances	253	84.6154 %
Incorrectly Classified Instances	46	15.3846 %
Kappa statistic	0.7692	
Mean absolute error	0.1021	
Root mean squared error	0.3158	
Relative absolute error	22.9815 %	
Root relative squared error	66.99 %	
Total Number of Instances	299	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.949	0.035	0.931	0.949	0.940	0.910	0.987	0.975	classical
	0.710	0.025	0.934	0.710	0.807	0.742	0.974	0.938	hiphop
	0.880	0.171	0.721	0.880	0.793	0.681	0.921	0.778	rock
Weighted Avg.	0.846	0.077	0.862	0.846	0.846	0.777	0.961	0.897	

=== Confusion Matrix ===

a b c <-- classified as
94 0 5 | a = classical

0 71 29 | b = hiphop

7 5 88 | c = rock

=====J48=====

Correctly Classified Instances	248	82.6667 %
Incorrectly Classified Instances	52	17.3333 %
Kappa statistic	0.74	
Mean absolute error	0.1225	
Root mean squared error	0.3324	
Relative absolute error	27.5682 %	
Root relative squared error	70.5044 %	
Total Number of Instances	300	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.880	0.045	0.907	0.880	0.893	0.842	0.923	0.846	classical
	0.850	0.085	0.833	0.850	0.842	0.761	0.895	0.763	hiphop
	0.750	0.130	0.743	0.750	0.746	0.618	0.796	0.654	rock
Weighted Avg.	0.827	0.087	0.828	0.827	0.827	0.740	0.871	0.754	

=== Confusion Matrix ===

a b c <-- classified as

88 1 11 | a = classical

0 85 15 | b = hiphop

9 16 75 | c = rock

=====SMO=====

Correctly Classified Instances	288	96 %
Incorrectly Classified Instances	12	4 %
Kappa statistic	0.94	
Mean absolute error	0.2311	
Root mean squared error	0.288	
Relative absolute error	52 %	
Root relative squared error	61.101 %	
Total Number of Instances	300	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.990	0.000	1.000	0.990	0.995	0.993	0.999	0.998	classical
	0.920	0.015	0.968	0.920	0.944	0.917	0.979	0.938	hiphop
	0.970	0.045	0.915	0.970	0.942	0.912	0.963	0.898	rock
Weighted Avg.	0.960	0.020	0.961	0.960	0.960	0.941	0.980	0.944	

=== Confusion Matrix ===

a b c <-- classified as
99 0 1 | a = classical
0 92 8 | b = hiphop
0 3 97 | c = rock

Using scikit-learn:

	Support Vector Machine	Stochastic Gradient Descent	Nearest Neighbours
Confusion matrix:	$\begin{bmatrix} 23 & 0 & 1 \\ 0 & 23 & 3 \\ 0 & 0 & 25 \end{bmatrix}$	$\begin{bmatrix} 24 & 0 & 0 \\ 3 & 23 & 0 \\ 7 & 18 & 0 \end{bmatrix}$	$\begin{bmatrix} 23 & 0 & 1 \\ 2 & 22 & 2 \\ 3 & 6 & 16 \end{bmatrix}$
Accuracy:	0.96666666666667	0.64	0.88666666666667

Question 2

▪

```
69 #start
70 tag = {'Rap' : 12, 'Pop_Rock' : 1, 'Country' : 3}
71 words = get_words()
72 tracks = get_tracks()
73 labels = genres()
```

```

75 def extract_vocabulary(D):
76     V = []
77     for word in words:
78         V.append(word)
79     return V
80
81 def count_documents(D):
82     return len(D)
83
84 def count_docs_in_class(D, c):
85     return len([get_tracks(genre = c)])
86
87 def concatenate_all_text_in_docs(D, c):
88     text, tracks = [], get_tracks(genre = c)
89     for word in words:
90         for track in tracks:
91             n = track.num_word(words.index(word))
92             while n > 0:
93                 text.append(word)
94                 n = n - 1
95     return text
96
97 def count_tokens_of_terms(text_c, t):
98     return text_c.count(t)
99
100 def extract_tokens_from_doc(V, d):
101     text = []
102     for word in words:
103         n = d.num_word(words.index(word))
104         while n > 0:
105             text.append(word)
106             n = n - 1
107     return text
108
109 #train the multinomial model
110 def train_multinomial(C, D):
111     V = extract_vocabulary(D)
112     N = count_documents(D)
113     prior, condprob = dict.fromkeys(C, {})
114     for c in C:
115         N_c = count_docs_in_class(D, c)
116         prior[c] = float(N_c)/N
117         text_c = concatenate_all_text_in_docs(D, c)
118         T = {}
119         for t in V:
120             T[(t,c)] = count_tokens_of_terms(text_c, t)
121         for t in V:
122             condprob[(t,c)] = float(T[(t,c)] + 1)
123                             /(sum([T[(tp,c)] + 1 for tp in V]))
124     return V, prior, condprob
125
126 #apply the multinomial on new instance d
127 def apply_multinomial(C, V, prior, condprob, d):
128     W = extract_tokens_from_doc(V, d)
129     score = dict.fromkeys(C, 0)
130     for c in C:
131         score[c] = np.log(prior[c])
132         for t in W:
133             score[c] += np.log(condprob[(t,c)])
134     argmax, max_score = '', -np.inf
135     for c in C:
136         if score[c] > max_score:
137             max_score = score[c]
138             argmax = c
139     return argmax

```

The probability of a word given a genre, $pr(\text{word} \mid \text{genre})$, is obtained by `condprob[(word,genre)]` which can be obtained by calling `train_multinomial(C, D)`. (lines 110 - 124 above) (Resulting table below)

(de,Rap,0.072)	(de,Pop_Rock,0.026)	(de,Country,0.002)
(niggaz,Rap,0.042)	(niggaz,Pop_Rock,0.003)	(niggaz,Country,0.001)
(ya,Rap,0.093)	(ya,Pop_Rock,0.012)	(ya,Country,0.012)
(und,Rap,0.047)	(und,Pop_Rock,0.019)	(und,Country,0.000)
(yall,Rap,0.042)	(yall,Pop_Rock,0.002)	(yall,Country,0.005)
(ich,Rap,0.067)	(ich,Pop_Rock,0.030)	(ich,Country,0.000)
(fuck,Rap,0.069)	(fuck,Pop_Rock,0.022)	(fuck,Country,0.002)
(shit,Rap,0.093)	(shit,Pop_Rock,0.006)	(shit,Country,0.003)
(yo,Rap,0.078)	(yo,Pop_Rock,0.009)	(yo,Country,0.005)
(bitch,Rap,0.059)	(bitch,Pop_Rock,0.004)	(bitch,Country,0.001)
(end,Rap,0.014)	(end,Pop_Rock,0.037)	(end,Country,0.022)
(wait,Rap,0.013)	(wait,Pop_Rock,0.045)	(wait,Country,0.026)
(again,Rap,0.019)	(again,Pop_Rock,0.048)	(again,Country,0.053)
(light,Rap,0.016)	(light,Pop_Rock,0.044)	(light,Country,0.032)
(eye,Rap,0.023)	(eye,Pop_Rock,0.056)	(eye,Country,0.042)
(noth,Rap,0.012)	(noth,Pop_Rock,0.038)	(noth,Country,0.021)
(lie,Rap,0.009)	(lie,Pop_Rock,0.038)	(lie,Country,0.017)
(fall,Rap,0.011)	(fall,Pop_Rock,0.050)	(fall,Country,0.031)
(our,Rap,0.023)	(our,Pop_Rock,0.062)	(our,Country,0.043)
(away,Rap,0.017)	(away,Pop_Rock,0.079)	(away,Country,0.054)
(gone,Rap,0.016)	(gone,Pop_Rock,0.035)	(gone,Country,0.044)
(good,Rap,0.029)	(good,Pop_Rock,0.033)	(good,Country,0.062)
(night,Rap,0.023)	(night,Pop_Rock,0.063)	(night,Country,0.071)
(blue,Rap,0.007)	(blue,Pop_Rock,0.015)	(blue,Country,0.037)
(home,Rap,0.015)	(home,Pop_Rock,0.034)	(home,Country,0.055)
(long,Rap,0.017)	(long,Pop_Rock,0.037)	(long,Country,0.065)
(liittl,Rap,0.025)	(liittl,Pop_Rock,0.038)	(liittl,Country,0.075)
(well,Rap,0.022)	(well,Pop_Rock,0.044)	(well,Country,0.065)
(heart,Rap,0.015)	(heart,Pop_Rock,0.052)	(heart,Country,0.087)
(old,Rap,0.011)	(old,Pop_Rock,0.019)	(old,Country,0.066)

More code used in Q2:

```
142 def run_multinomial():
143
144     C, D, k = ['Rap', 'Pop_Rock', 'Country'], get_tracks(), 10
145
146     #applies the multinomial and prints the data
147     def print_statistics(data):
148         confusion_matrix = [[0,0,0,C[2]], [0,0,0,C[1]], [0,0,0,C[0]]]
149         total = 0
150         for d in [d for d in data if d.genre == tag['Rap']]:
151             c = apply_multinomial(C, V, prior, condprob, d)
152             if c == C[0]:
153                 total += 1
154                 confusion_matrix[2][0] += 1
155             elif c == C[1]:
156                 confusion_matrix[2][1] += 1
157             elif c == C[2]:
158                 confusion_matrix[2][2] += 1
159
160         for d in [d for d in data if d.genre == tag['Pop_Rock']]:
161             c = apply_multinomial(C, V, prior, condprob, d)
162             if c == C[0]:
163                 confusion_matrix[1][0] += 1
164             elif c == C[1]:
165                 total += 1
166                 confusion_matrix[1][1] += 1
167             elif c == C[2]:
168                 confusion_matrix[1][2] += 1
169
170         for d in [d for d in data if d.genre == tag['Country']]:
171             c = apply_multinomial(C, V, prior, condprob, d)
172             if c == C[0]:
173                 confusion_matrix[0][0] += 1
174             elif c == C[1]:
175                 confusion_matrix[0][1] += 1
176             elif c == C[2]:
177                 total += 1
178                 confusion_matrix[0][2] += 1
179
180         print("classification accuracy: " + str(total*1.0/count_documents(data)))
181         print('|%10s |%10s |%10s |' % (C[0],C[1],C[2]))
182         print('-----')
183         for row in confusion_matrix:
184             for col in row:
185                 print('|%10s' % col),
186             print('')
187         print('\n')
188
189     #test on kth fold and train on the remaining k-1 folds for each k
190     random.shuffle(D)
191     V, prior, condprob = train_multinomial(C, D)
192     for i in range(k):
193         k_folds = np.split(D, k)
194         test_data = k_folds.pop(i)
195         train_data = [j for i in k_folds for j in i]
196         V, prior, condprob = train_multinomial(C, train_data)
197         print_statistics(test_data)
```

Confusion Matrix for Naive Bayes classifier.

Using all sets as training data and test data:

classification accuracy: 0.68

	Rap		Pop_Rock		Country	

	25	206	769	Country
	114	507	379	Pop_Rock
	748	118	134	Rap

Using k = 10 folds, use 9 sets for training and 1 for testing. Run 10 times to try each kth fold as test data with the remaining k-1 folds as training data:

classification accuracy: 0.68	classification accuracy: 0.71
Rap Pop_Rock Country	Rap Pop_Rock Country
-----	-----
3 23 70 Country	3 18 81 Country
11 56 33 Pop_Rock	7 53 37 Pop_Rock
78 15 11 Rap	79 10 12 Rap
classification accuracy: 0.683333333333	classification accuracy: 0.653333333333
Rap Pop_Rock Country	Rap Pop_Rock Country
-----	-----
4 23 84 Country	2 23 81 Country
10 46 36 Pop_Rock	15 52 37 Pop_Rock
75 10 12 Rap	63 12 15 Rap
classification accuracy: 0.62	classification accuracy: 0.666666666667
Rap Pop_Rock Country	Rap Pop_Rock Country
-----	-----
2 20 58 Country	1 13 77 Country
15 52 46 Pop_Rock	7 46 46 Pop_Rock
76 15 16 Rap	77 18 15 Rap
classification accuracy: 0.68	classification accuracy: 0.693333333333
Rap Pop_Rock Country	Rap Pop_Rock Country
-----	-----
2 23 86 Country	2 23 72 Country
13 48 38 Pop_Rock	13 64 33 Pop_Rock
70 8 12 Rap	72 11 10 Rap
classification accuracy: 0.67	classification accuracy: 0.69
Rap Pop_Rock Country	Rap Pop_Rock Country
-----	-----
2 21 77 Country	4 19 83 Country
14 44 37 Pop_Rock	9 46 36 Pop_Rock
80 13 12 Rap	78 6 19 Rap

Making randomly generated tracks using the probability distribution of a word occurring given a genre:


```

208     def get_probabilistic_word(genre = None):
209         assert(genre is not None)
210         prob_dist = []
211         for word in words:
212             prob_dist.append(condprob[(word, genre)])
213         return np.random.choice(words, p = prob_dist)
214     n_lyrics, n_songs = 20, 5
215     generated_tracks = []
216     for e, c in enumerate(C):
217         for i in range(n_songs):
218             t = Track(n_songs*e+i, [0]*30, c)
219             for j in range(n_lyrics):
220                 t.add_word(words.index(get_probabilistic_word(c)))
221             generated_tracks.append(t)
222     for t in generated_tracks:
223         t.print_track()

```

ID: 0 Genre: 12

Feature Vector: [0, 2, 0, 2, 0, 3, 2, 2, 1, 1, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 2, 0]

ID: 1 Genre: 12

Feature Vector: [0, 2, 1, 2, 2, 1, 3, 5, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0]

ID: 2 Genre: 12

Feature Vector: [1, 1, 0, 0, 2, 1, 1, 4, 3, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0]

ID: 3 Genre: 12

Feature Vector: [3, 0, 1, 2, 1, 0, 3, 2, 1, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 2, 0, 0, 0, 0]

ID: 4 Genre: 12

Feature Vector: [2, 2, 1, 1, 2, 0, 1, 1, 2, 2, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0]

ID: 5 Genre: 1

Feature Vector: [0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 2, 0, 2, 2, 1, 0, 3, 1, 1, 1, 2, 0, 0, 1, 0, 0, 1, 0]

ID: 6 Genre: 1

Feature Vector: [0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 2, 1, 2, 3, 1, 0, 0, 3, 0, 2, 0, 1, 2, 1, 0, 0, 0, 0]

ID: 7 Genre: 1

Feature Vector: [0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 2, 1, 0, 1, 0, 0, 2, 2, 1, 0, 0, 1, 0, 1, 1, 2, 2, 0, 3]

ID: 8 Genre: 1

Feature Vector: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 1, 0, 1, 0, 0, 2, 3, 2, 1, 1, 3, 1, 2, 0, 0, 0, 0, 1]

ID: 9 Genre: 1

Feature Vector: [0, 0, 0, 1, 0, 2, 0, 1, 1, 0, 1, 0, 2, 0, 0, 0, 2, 1, 0, 1, 1, 1, 1, 0, 2, 0, 1, 1, 1, 0]

ID: 10 Genre: 3

Feature Vector: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 3, 2, 0, 0, 1, 2, 0, 1, 1, 0, 0, 0, 0, 1, 2, 4, 0]

ID: 11 Genre: 3

Feature Vector: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 2, 1, 0, 0, 0, 1, 2, 0, 2, 1, 2, 2, 0, 0, 1, 3, 1, 1]

ID: 12 Genre: 3

Feature Vector: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 3, 1, 3, 0, 0, 0, 1, 3, 0, 0, 3, 0, 0, 0, 1, 1, 0, 4]

ID: 13 Genre: 3

Feature Vector: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 2, 0, 0, 1, 2, 0, 1, 2, 2, 0, 4, 3]

ID: 14 Genre: 3

Feature Vector: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 2, 2, 1, 1, 1, 0, 1, 2, 2, 3, 1]

Running classifier on generated data for fun :)

classification accuracy: 0.866666666667

Rap	Pop_Rock	Country
-----	----------	---------

0	0	5	Country
0	3	2	Pop_Rock
5	0	0	Rap

Happy Marking!